

# Fast and Furious Withdrawals from Optimistic Rollups

**Abstract.** Optimistic rollups are in wide use today as a scalability add-on for blockchains like Ethereum. In such systems, Ethereum is referred to as L1 (Layer 1) and the rollup provides an environment called L2, which reduces fees and latency but cannot instantly and trustlessly interact with L1. One practical issue for optimistic rollups is that trustless transfers of tokens and ETH, as well as general messaging, from L2 to L1 is not finalized on L1 until the passing of a dispute period (aka withdrawal window) which is currently 7 days in the two leading optimistic rollups: *Arbitrum* and *Optimism*. In this paper, we explore methods for side-stepping the dispute period when withdrawing ETH from L2. We modify a popular rollup, *Arbitrum*, to enable withdraws to be traded on L1 before they are finalized, and combine them with a prediction market on L1 that effectively provides insurance in the event that the withdraws do not finalize. As a result, anyone (including contracts) on L1 can safely accept withdrawn tokens while the dispute period is open despite having no knowledge of what is happening on L2. All fees are set by open market operations.

**Keywords:** Ethereum · layer 2 · rollups · bridges · prediction markets

## 1 Introductory Remarks

Ethereum-compatible blockchain environments, called Layer 2s (or L2s), have demonstrated an ability to reduce transaction fees by 99–99.9% while preserving the strong guarantees of integrity and availability in the underlying blockchain. The subject of this paper concerns one sub-category of L2 technology called an optimistic rollup. The website *L2 Beat* attempts to capitalize all tokens of known value across the top 25 L2 projects, and finds that two optimistic rollups, *Arbitrum* and *Optimism*, respectively account for 50% and 30% of L2 value—\$4B USD at the time of writing.<sup>1</sup>

We will describe the working details of optimistic rollups later in this paper but here are the main takeaways. Currently, rollups are faster and cheaper than Ethereum itself, however each L2 is essentially an isolated environment that cannot instantly and trustlessly interact with accounts and contracts that are running on either L1 or other L2s. An optimistic rollup project will typically provide a smart contract, called a validating bridge [5], that can trustlessly move ETH (and other tokens and even arbitrary messages) between L1 and its own

---

<sup>1</sup> L2 Beat, accessed Oct. 2022.

L2. It does this by locking the ETH in the L1 contract and later releasing it, on request, according to who its new owner is on L2 at the time of the request. If finalized, the ETH will be destroyed on L2 and will be released by the bridge on L1.

Owing to how optimistic rollups convince the L1 bridge contract of who the current owner of withdrawn ETH is on L2 (explained below), the bridge has to first wait a period of time called the dispute window. The current default is 7 days in *Arbitrum* and *Optimism*, however the filing of new disputes can extend the window. The bottom line is that users have to wait at least 7 days to draw down ETH from an optimistic rollup.

*Contributions.* In this paper, we compare several methods—atomic swaps and tradable exits—for working around this limitation. While we argue workarounds cannot be done generally, some circumstances allow it: namely, when the withdrawn token is liquid, fungible, and available on L1 and the withdrawer is willing to pay a fee to speed up the withdraw. We concentrate in the task of drawing down ETH specifically but the solution works for any fungible ERC20 (or related) token. While these techniques work easily between human participants that have off-chain knowledge, such as the valid state of the L2, it is harder to make them compatible with L1 smart contracts that have no ability to validate the state of L2. We propose a solution using tradable exits and prediction markets to enable an L1 smart contract to safely accept withdrawn tokens before the dispute period is over. Finally, we modify the current version, *Nitro*, of the most popular optimistic rollup, *Arbitrum*, made open source<sup>2</sup> by *Offchain Labs* to implement our solution and provide measurements. *Arbitrum* is a commercial product with academic origins [4].

## 2 Background

*Inbox.* Roll-ups have emerged as a workable approach to reducing fees and latency for Ethereum-based decentralized applications. In a roll-up, transactions to be executed on L2 are recorded in an L1 smart contract called the inbox. Depending on the system, users might submit to the inbox directly, or they might submit to an offchain service, called a sequencer, that will batch together transactions from many users and pay the L1 fees for depositing them in the inbox. Transactions recorded in the inbox (as `calldata`) are not executed on Ethereum. Instead, they are executed in a separate environment off the Ethereum chain, called L2. This external environment will operate by different rules designed to reduce fees, increase throughput, and decrease latency.

*Outbox.* Occasionally (*e.g.*, every 5 minutes), validators on L2 will produce a checkpoint of the state of all contracts and accounts in the complete L2 according to the latest transactions and will place this asserted state in a contract on L1 called the outbox. Note that anyone with a view of L1 can validate that the

---

<sup>2</sup> GitHub: Nitro

sequence of transactions recorded in the inbox produces the asserted checkpoint in the outbox, however asking Ethereum to validate this be equivalent to running the transactions on Ethereum. The key breakthrough is that the assertion will be posted with *evidence* that the checkpoint is correct.

*Optimistic vs. zk-rollups.* In practice, two main types of evidence are used. In zk-rollups,<sup>3</sup> a succinct computational argument that the assertion is correct is posted and can be checked by Ethereum for far less cost than running all of the transactions. However the proof is expensive to produce. In optimistic rollups, the assertions are backed by a large amount of cryptocurrency (acting as a fidelity bond). The correctness of the assertion can be disputed by anyone on Ethereum and Ethereum itself can decide between two (or more) disputes for far less cost than running all of the transactions. It will then reallocate the fidelity bonds to the whoever is making correct assertions. If an assertion is undisputed for a window of time (*e.g.*, 7 days), the assertion is considered final.

*Bridge.* A final piece of layer 2 infrastructure is a bridge, which can move ETH, tokens, NFTs, and even arbitrary messages, between layer 1 and layer 2. In this paper, we consider the case of a bridge for ETH (while discussing other kinds much later in Section ??). If Alice has ETH on Ethereum, she can submit her ETH to a bridge smart contract on Ethereum which will lock the ETH inside of it, while generating the same amount of ETH in Alice’s account inside the layer 2 environment. Since any deviation from this will result in an incorrect checkpoint which is ultimately checkable by Ethereum, the bridge does not need to be trusted. A set of transactions on layer 2 might see the ETH move from Alice’s account to Bob’s account on Layer 2. Bob is now entitled to draw down from layer 2 to layer 1. He submits a withdraw request, which reduces his Layer 2 balance and ends up in the next asserted checkpoint on Layer 1. Optimistically, the checkpoint is undisputed and 7 days later it is finalized. Bob can now ask the bridge to release the ETH by demonstrating his withdraw is contained in the finalized checkpoint.

## 2.1 Related Work

Arbitrum is described at *USENIX Security* [4]. Gudgeon *et al.* provide a systemization of knowledge (SoK) of Layer 2 technology (that largely predates rollups) [3], while McCorry *et al.* provide an SoK that covers roll-ups and validating bridges [5]. Some papers implement research solutions on Arbitrum for improved performance: decentralized order books [6] and secure multiparty computation [2]. Further academic work on optimistic rollups and bridges is largely missing at this time, but we anticipate it will become an important research area. Other related topics are atomic swaps and prediction markets. Too many papers propose atomic swap protocols to list here but see Zamyatin *et al.* for an SoK

<sup>3</sup> zk stands for zero-knowledge, a slight misnomer: succinct arguments of knowledge that only need to be complete and sound, not zero-knowledge, are used.

of the area (and a new theoretical result) [8]. Decentralized prediction markets proposals predate Ethereum and include Clark *et al.* [1] and Truthcoin [7]. Early Ethereum projects Augur and Gnosis began as prediction markets.

### 3 Proposed Solution

- Optimistic rollups are being used
- Problem: multi-day (*e.g.*, seven) window for withdrawal
- Why is this a problem? Others are solving this problem (alternative bridges with a TTP). Some examples include: Speculators want to move fast, voting in a DAO (red tape), sell them on L1 or another L2, DApp access on L1 (trading is efficient on L2 but you need to do something on L1).
- Disadvantage relative to zk-rollups
- Solution: scope to liquid tokens (ETH and ERC20), open problem: NFTs or other non-substitutable tokens.
- Solution: (1) tradeable exits; (2) market to trade; (3) guaranteed exit (buyer runs ArbOS validator); (4) prediction market solution to guarantee exits to non-validating entities (importantly includes smart contracts)
- Testing: we implemented (1) and (3); for (2) and (4), use your favourite DeFi project.

#### 3.1 Design Landscape

*Notation.* Consider an amount of 100 ETH. When this amount is in the user's account on L1, we use the notation 100 ETH<sub>L1</sub>. When it is in the bridge on L1 and in the user's account on L2, we denote it 100 ETH<sub>L2</sub>. When the ETH has been withdrawn on L2 and the withdrawal has been asserted in the L1 outbox, but the dispute window is still open, we refer to it as 100 ETH<sub>xx</sub>. Other transitional states are possible but not needed for our purposes.

- Landscape: atomic swaps
- Landscape: Change the bridge instead of using a third party contract (reason: too late once you withdraw)

#### 3.2 Fast Bridge

- Allow trading of exits (atomic unit) -> track most recent owner (constant time) in outbox. Authorization (only current owner can transfer) -> on execute, check for current owner.
- Self-insurance (staking a fidelity bond) (you need liquidity)
- Prediction market: someone else will insure it (better than insurance because you can exit quickly, change your mind, make money on over-/under- without necessarily agreeing with the position you are buying)

### 3.3 Prediction Market

In our protocol so far, Alice wants to fast withdraw 100  $\text{ETH}_{/L2}$ . Bob has 100  $\text{ETH}_{/L1}$  that he will not use until after the dispute window. Bob also runs an L2 validator so he is assured that if Alice withdraws, it is valid and will eventually finalize. In this case, Alice will withdraw 100  $\text{ETH}_{/XX}$  and swap it for Bob's 100  $\text{ETH}_{/L1}$ , while paying a fee to Bob.

One remaining issue with this method is how specialized Bob is: he must have liquidity in  $\text{ETH}_{/L1}$ , be an active trader who knows how to price futures, and be an L2 validator. While we can expect blockchain participants with each specialization, it is a lot to assume of a single entity. The goal of this subsection is to split Bob into two distinct participants: one that has  $\text{ETH}_{/L1}$  liquidity but does not know about L2 (Carol) and one that knows about L2 but is not necessarily an active trader on L1 (David). The main impact of this change is that Carol can be an autonomous L1 smart contract. Carol might be an intermediary that exchanges  $\text{ETH}_{/XX}$  and  $\text{ETH}_{/L1}$ . Alternatively, recall that Alice wants  $\text{ETH}_{/L1}$  quickly in order to do something on L1 with it; Carol can be that destination account or contract.

What is the risk if Carol just accepts  $\text{ETH}_{/XX}$  as if it were  $\text{ETH}_{/L1}$ ? Carol can check, using L1 only, that Alice owns  $\text{ETH}_{/XX}$  within a disputable assertion that is registered in the L1 outbox of the rollup. The risk is that the assertion is wrong and is not finalized.

### 3.4 Implementation

- Arbitrum's Nitro. Bridge: inbox, sequencer, outbox.
- Implemented a market, following is the gas cost related to the market:  
gasUsed for opening a market on an exit: 328,029 gasUsed for transferring the exit to the market: 86,701 (it was the first transfer so a bit more expensive than the 2nd, 3rd,.. see below) Gas cost for submitting the Bid for when the Bid is greater than ask -> trade occurs and settles in one single submitBid tx: 105,287 Gas cost for execution is: 92,148
- Modify outbox to allow tradeable exits
- Modify the Nitro codebase (arbnode and validator) to shrink the fraud proof window from 7 days to 1 minute: (1) Modify the confirmPeriod variable in the arbnode/node.go file, (2) modify the MakeAssertionInterval variable validator/staker.go
- To make the bridge prediction market friendly: Modify the outbox: (1) added a Mapping that maps the proposed arbitrum block number (also known as assertion and node) to the pending state. (2) added a function which accepts a proposed block number and adds it to the pending assertions mapping. Modify the rollupcore: The validator acts through the Rollup-Core.sol contract when making an assertion by calling a createNewNode() function. We modifies this function so that every time a node is created by the validator it's also added to the outbox's pending assertions mapping (outbox.addToPendingAssertions(latestNodeCreated()))

- L1 gas costs: new function (transferSpender) : First Transfer: 1) Alice withdraws ETH from L2 2) She transfers her exit to Bob transferSpender in this case costs : gasUsed : 85,945 Second Transfer: 2) Now Bob transfer his exit to Carol transferSpender in this case costs : gasUsed : 48,810 Third Transfer: 2) Now Carol transfer his exit to Nancy transferSpender in this case costs : gasUsed : 48,798 (difference of two mappings)
- L1 gas costs: execute the exit: 91,418
- Unit tests: say something
- What happens when an assertion fails? (pro: ticket fails with assertion, better for prediction markets (betting on assertion which is a batch of exits); ticket passes even if assertion fails, sounds better (caveat: probably won't happen))
- Challenges: SDK, where to change, unit test failed assertions (good assertion, bad assertion where withdrawal is ok, bad assertion where the withdrawal is problematic)
- Expose assertion failures/successes to external contract (submit ID for assertion, get back status: pending, finalized, or discarded). Write down the SDK call. (what happens to a failed assertions???)

## 4 Discussion

### 4.1 Where to Solve Problem?

- Exit to a third party collateral contract, implements trading

### 4.2 Withdrawal Format

- Divisible exits
- ERC20 / ERC721? -> allowances or not?

### 4.3 Non-Substitutable Withdrawals

- Illiquid tokens
- NFTs (substitute for ETH -> support)
- Messages (oracle updates, read calls, any L2-to-L1 message...) (still offer insurance)

### 4.4 Assertion Failures

- does exit go away with the assertion or not? Pros/cons

### 4.5 Markets

- Can't use an AMM
- Can't run out
- Simple auction
- Divisible exits
- Exit pools

## 4.6 Pricing

How much is 100  $\text{ETH}_{\text{XX}}$  worth in  $\text{ETH}_{\text{L1}}$ ? Using intuition, we have been treating as worth something like 99  $\text{ETH}_{\text{L1}}$  but we now describe the factors that impact the price ( $F_0$ ). We observe that  $\text{ETH}_{\text{XX}}$  is like a futures contract: it is the delivery of ETH in a certain number of days. Here is a standard set of terms used to price a futures contract, and the translation into our setting:

- *Spot price ( $S_0$ )*. 100  $\text{ETH}_{\text{XX}}$  is equivalent to the delivery of 100  $\text{ETH}_{\text{L1}}$  at some future settlement date. The value of 100  $\text{ETH}_{\text{L1}}$  today is 100  $\text{ETH}_{\text{L1}}$ .
- *Exchange risk*. We assume that the currency or token used to purchase an exit is the same as the currency or token that will be released by the outbox; thus no exchange risk.
- *Settlement time ( $\Delta t$ )*. The time until the  $\text{ETH}_{\text{L1}}$  can be withdrawn. In reality, this time is variable, depending on dispute. We simplify and assume it is known (*e.g.*, always 7 days from the assertion time). As  $\Delta t$  approaches 0, the  $F_0$
- *Storage cost ( $U$ )*. Securing  $\text{ETH}_{\text{XX}}$  and securing  $\text{ETH}_{\text{L1}}$  is identical, so not having to take possession of  $\text{ETH}_{\text{L1}}$  for  $\Delta t$  time does not reduce costs for a  $\text{ETH}_{\text{XX}}$  holder (the way it might for a commodity that has to be stored).
- *Cashflow / Yield ( $r$ )* It is theoretically possible to use  $\text{ETH}_{\text{L1}}$  to earn interest from underlying
- *Delivery cost ( $D$ )*. The  $\text{ETH}_{\text{XX}}$  holder needs to pay the gas fees for retrieving  $\text{ETH}_{\text{L1}}$ . This is inverted from a typical futures contract with physical delivery where the holder of the futures does not pay.
- *Settlement risk ( $k$ )*.

$$F_0 = (S_0 + U - D) \cdot e^{(r-q-k) \cdot \Delta t}$$

## References

1. Clark, J., Bonneau, J., Felten, E.W., Kroll, J.A., Miller, A., Narayanan, A.: On decentralizing prediction markets and order books. In: Workshop on the Economics of Information Security, State College, Pennsylvania. vol. 188 (2014)
2. Demirag, D., Clark, J.: Absentia: Secure multiparty computation on ethereum. In: International Conference on Financial Cryptography and Data Security. pp. 381–396. Springer (2021)
3. Gudgeon, L., Moreno-Sanchez, P., Roos, S., McCorry, P., Gervais, A.: Sok: Off the chain transactions. IACR Cryptol. ePrint Arch. **2019**, 360 (2019)
4. Kalodner, H., Goldfeder, S., Chen, X., Weinberg, S.M., Felten, E.W.: Arbitrum: Scalable, private smart contracts. In: USENIX Security Symposium. pp. 1353–1370 (2018)
5. McCorry, P., Buckland, C., Yee, B., Song, D.: Sok: Validating bridges as a scaling solution for blockchains. Cryptology ePrint Archive (2021)
6. Moosavi, M., Clark, J.: Lissy: Experimenting with on-chain order books. arXiv preprint arXiv:2101.06291 (2021)

7. Sztorc, P.: Truthcoin. peer-to-peer oracle system and prediction marketplace. (2015)
8. Zamyatin, A., Al-Bassam, M., Zindros, D., Kokoris-Kogias, E., Moreno-Sanchez, P., Kiayias, A., Knottenbelt, W.J.: Sok: Communication across distributed ledgers. In: International Conference on Financial Cryptography and Data Security. pp. 3–36. Springer (2021)